




Training loss (Eq. 1)

$$\mathcal{L}_{NLL} = -\frac{1}{N} \sum_i \left[y_i \log \frac{\alpha_i}{\alpha_i + \beta_i} + (1 - y_i) \log \frac{\beta_i}{\alpha_i + \beta_i} \right]$$

Beta-Bernoulli NLL; Adam, lr 1e-3, batch 32, 50 epochs

KL regulariser (Eq. 2) evaluated only as an ablation